


What the agent sees and does



One screenshot per step (1360 × 768). Nothing else: no page source, no accessibility tree, no hints.

Actions

click, double click, drag, scroll, type, press key, wait, done

Budget


Each task has a fixed step and time limit (for example 150 steps or 30 minutes). The same limit is used for both languages and for every model.

Instructions


Frozen. A test fails if any instruction changes, and numbers, IDs and quoted text are the same in every language.

The two machines

English VM



Arabic VM



Both machines start from the same image and the same files. The only difference is the system language: the Arabic machine has an Arabic interface, keyboard and folder names, laid out right to left.

How we grade

Recording

Every screenshot, every action and every network request is saved for each run.

Real websites, no side effects

For payments, applications and posts, a proxy captures the final request and does not send it. The captured request is the evidence.

Checker

One per task. It looks at the saved files, the application state or the captured request and returns one of three results. No LLM judge.

PASS

FAIL

EVAL_ERROR

EVAL_ERROR means something on our side broke (setup, recorder, checker, API). Those runs are thrown out and rerun. They are never counted as a model failure.

Pairs

We compare each task in English against the same task in Arabic. When one language passes and the other fails, that pair is rerun 3 times and checked by hand before it counts.