

What the agent sees and does



One screenshot per step

1360 × 768. No page source, no accessibility tree, no hints.



Mouse

click, double click, drag, scroll



Keyboard

type, press key




Control

wait, done

Budget

Each task has a fixed step and time limit (for example 150 steps or 30 minutes). The same limit is used for both languages and every model.

The two machines



*Both machines start from the same image and the same files.
Only the system language differs.*



Recording

Every screenshot, every action and every network request is saved for each run.



Real websites, no side effects

For payments, applications and posts, a proxy captures the final request and does not send it. The captured request is the evidence.



Frozen instructions

A test fails if any instruction changes. Numbers, IDs and quoted text are the same in every language.

How we grade



Checker

One per task. It looks at the saved files, the application state or the captured request. No LLM judge.

PASS

FAIL

EVAL_ERROR

EVAL_ERROR means something on our side broke (setup, recorder, checker, API). Those runs are thrown out and rerun. They are never counted as a model failure.



Pairs

Each task in English is compared with the same task in Arabic. When one language passes and the other fails, the pair is rerun 3 times and checked by hand before it counts.



Audit

Every pass is checked against its screenshots. Every fail gets a reason.