

Observation & action



One screenshot per step

1360 × 768 · no DOM, no accessibility tree, no hints



Mouse

click · double-click · drag · scroll



Keyboard

type · key (chords work under the Arabic layout)




Control

wait · done

Budget

Authorized per task (e.g. 150 steps / 30 min), frozen with the task, identical across languages and models.

Matched environments



Same disk image, same apps, same start state — only the locale differs.



Recorder

every screenshot, every action, every network request — kept as the audit bundle



Capture-and-block proxy

the final consequential request (payment, application, post) is captured as evidence and answered with a synthetic response — it never reaches the site



Frozen instructions

test-locked; numbers, IDs and quoted strings identical across all arms

Evaluation



Deterministic checker

One per task. Reads only the recorded evidence: files, application state, captured requests. No LLM judge in the primary metric.

PASS

FAIL

EVAL_ERROR

EVAL_ERROR = our fault (setup, recorder, checker, provider). Excluded and rerun — never counted as a model failure.



Paired cells

The unit is the pair: one task, two arms. Divergent pairs (one arm passes) are re-run $k = 3$ and adjudicated from evidence before they count.



Audit

Every PASS read against its screenshots; every FAIL labeled with a cause; per-family significance by paired McNemar test.