

Phase 1 • Before any run



1. Frozen instructions

Every English and Arabic instruction is locked by a test that compares it character-for-character to the approved source. Pinned literals must match across all six arms.



2. Build with truth-table tests

Setups and checkers are implemented against a written work order. Each checker ships with a passing fixture, a near-miss that must fail, and both language arms.



3. Independent verification gate

Nothing merges until a second reviewer checks the branch against the work order: scope, no invented matchers, no weakened tests, no secrets.



Phase 2 • During and after runs



4. Watched cells

One task per family runs end to end in both languages while a human reads the evidence. A live block-proof must show a real final click captured and blocked before consequential tasks run.



5. Batches under the EVAL_ERROR rule

Missing setup, recorder, or checker evidence is an evaluation error, never a zero. Environment defects are fixed as their own commits and the affected cells rerun fresh — results are never edited.



6. Evidence adjudication

Every PASS is read against its screenshots and captures; every FAIL gets a cause label; divergent pairs are re-run $k = 3$. A defect ledger records every fault found and how it was closed.